



# Network Socket Programming - II

---

BUPT/QMUL

2021-03-25



北京邮电大学

BEIJING UNIVERSITY OF POSTS AND TELECOMMUNICATIONS

Electronic Engineering 



# Review

---

- Basic Concepts in network programming
  - Process
  - File descriptor
  - System call
  - ~~Signal~~
- Some Helpful Points
  - Client-Server Model
  - Connections



# Agenda

---

- Week 2
  - Basic Concepts in Network Programming
  - Review of Some Helpful Points
- Week 4
  - Structures About IP Address and DNS
  - Sockets Interface
  - Major System Calls
  - Sample Programs



---

# Structures about IP address and DNS



# A programmer's view of the Internet

---

- 1. Hosts are mapped to a set of 32-bit *IP addresses*.
  - 202.112.96.163
- 2. The set of IP addresses is mapped to a set of identifiers called Internet *domain names*.
  - 202.112.96.163 is mapped to www.bupt.edu.cn
- 3. A process on one Internet host can communicate with a process on another Internet host over a *connection*.

# IP Addresses (1)

- 32-bit IP addresses are stored in an *IP Address structure*


- `<netinet/in.h>`

```
/* Internet address. */  
typedef uint32_t in_addr_t;  
struct in_addr {  
    in_addr_t s_addr;  
};
```

```
/*Defined in <stdint.h>*/  
typedef unsigned int uint32_t;
```

- Two ways to store multi-byte integers

- Big-endian vs. little-endian





# IP Addresses (2):

## Host byte order vs. network byte order

---

- **Host byte order** is machine-dependent
  - You can see it in `<bits/endian.h>`
  - A program used to output the host byte order
- **Network byte order** is machine-independent (*big-endian*)
- Byte order conversion functions
  - **htonl**: host byte order → network byte order for *long int*
  - **htons**: host byte order → network byte order for *short int*
  - **ntohl**: network byte order → host byte order for *long int*
  - **ntohs**: network byte order → host byte order for *short int*

# IP Addresses (3)

-- A program used to output the host byte order

## ■ byteorder.c

```
#include <sys/types.h>
#include <stdio.h>
#include <stdlib.h>

int main(void)
{
    union {
        short s;
        char c[sizeof(short)];
    } un;

    un.s=0x0102;

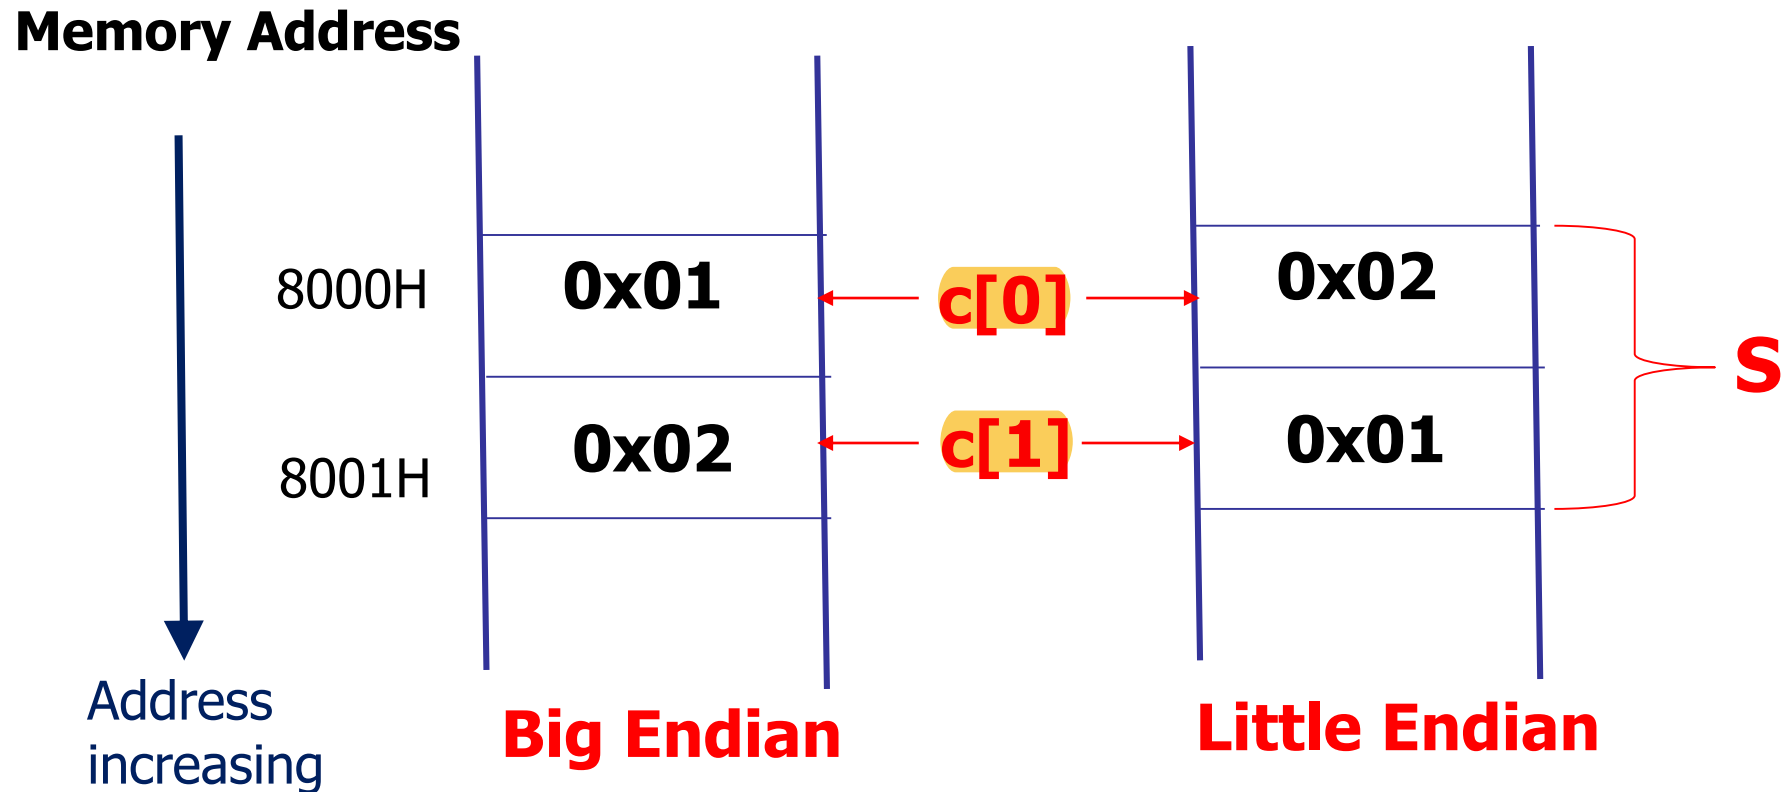
    if (sizeof(short)==2) {
        if (un.c[0]==1 && un.c[1]==2)
            printf("big-endian\n");
        else if (un.c[0]==2 && un.c[1]==1)
            printf("little-endian\n");
        else
            printf("Unknow\n");
    } else
        printf("sizeof(short)=%d\n", sizeof(short));

    exit(0);
}
```

a **union** is like a **struct**, only all the data members sit at the same memory location. This means only one of them can be used at a time.



# Storage example of Variable *s* and *c*





# Domain Name System (1)

---

- The Internet maintains a mapping between IP addresses and domain names in a huge worldwide distributed database called *DNS*.
  - Conceptually, programmers can view the DNS database as a collection of millions of *host entry structures*
  - `<netdb.h>`

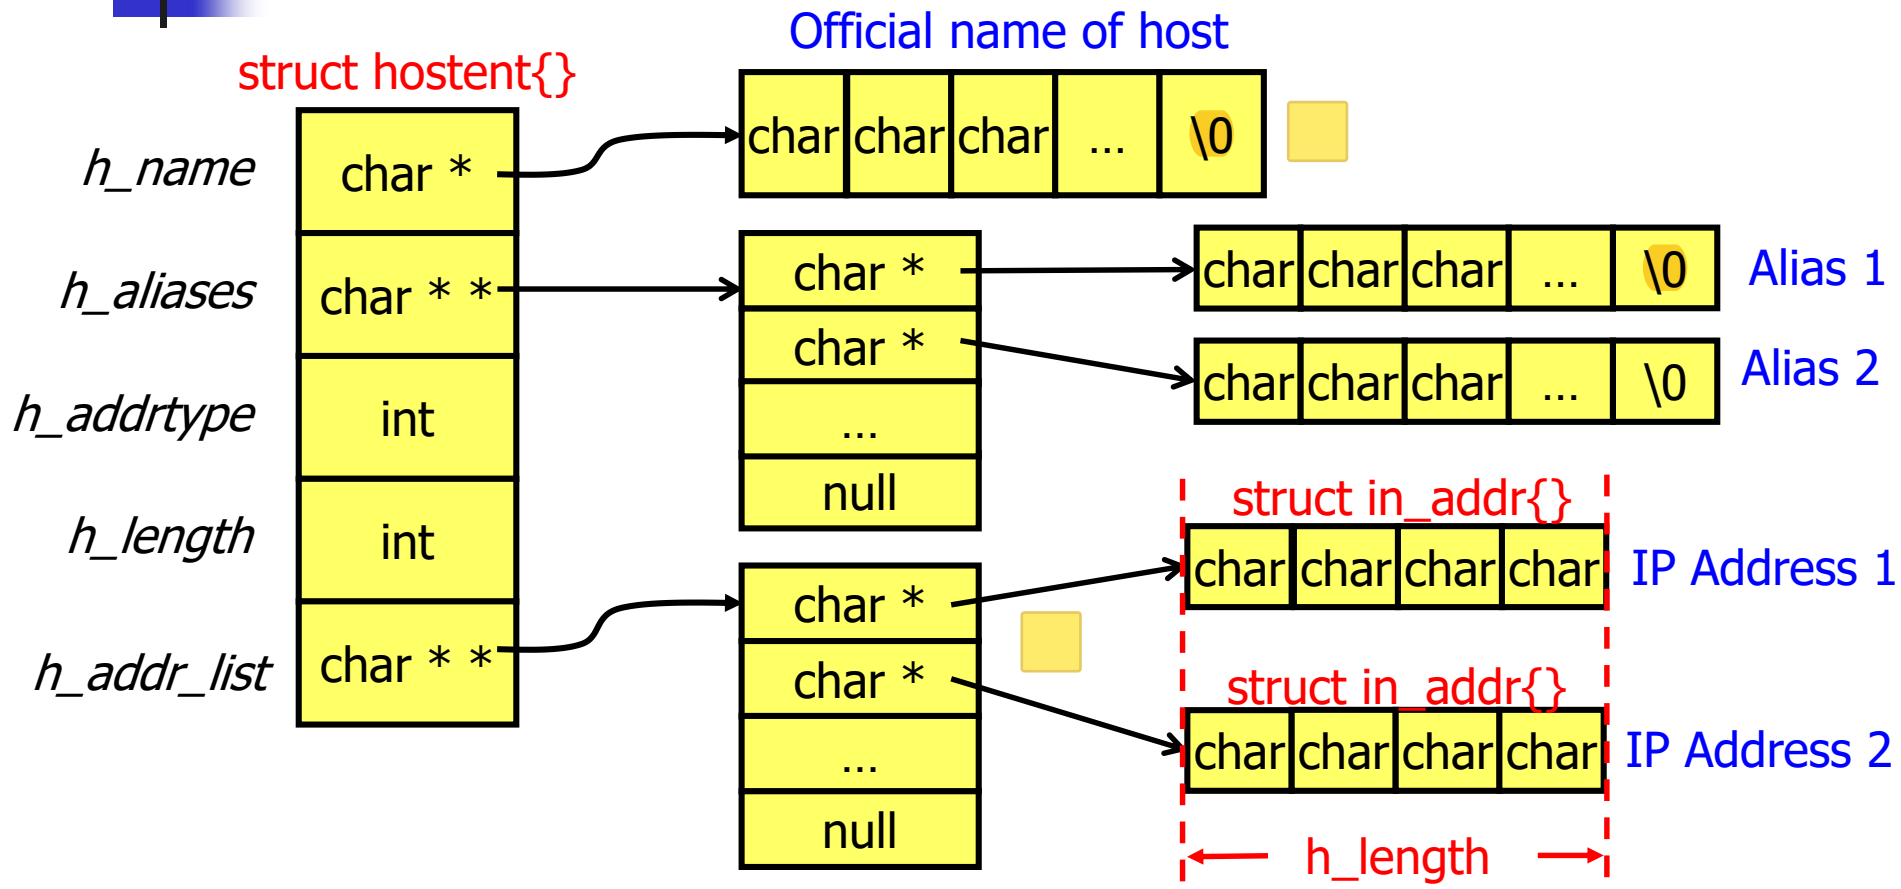
# Domain Name System (2): Host Entry Structure

## **h\_addr\_list:**

An array of pointers to IP addresses for the host (in network byte order), terminated by a NULL pointer.

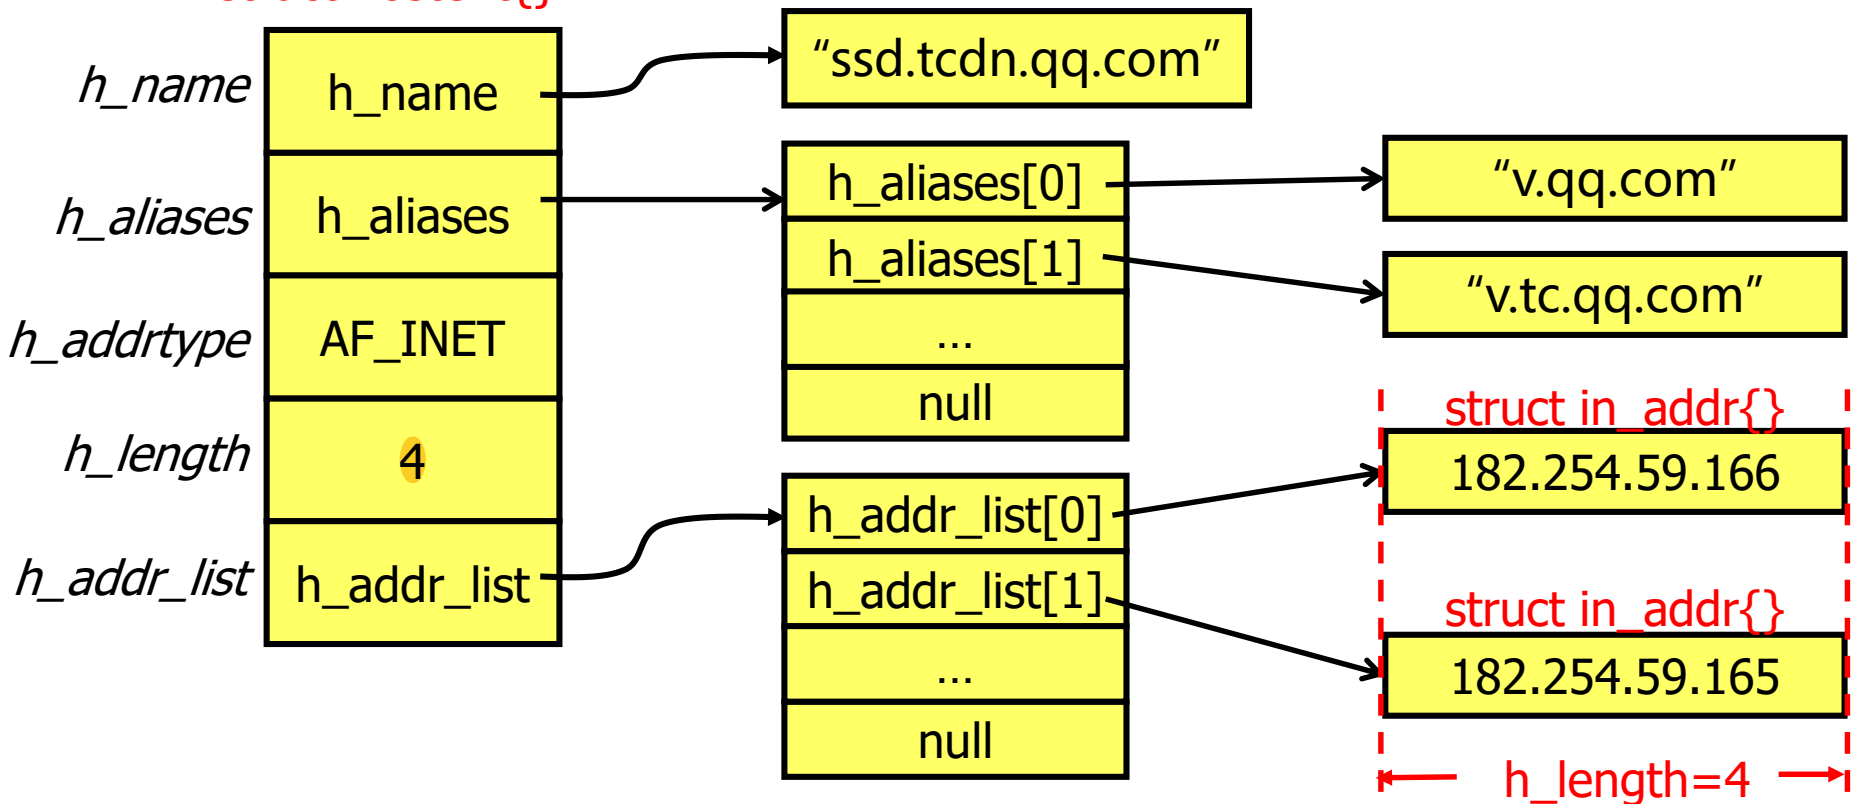
```
/* Description of data base entry for a single host. */
struct hostent
{
    char *h_name; /* Official name of host. */
    char **h_aliases; /* Alias list. */
    int h_addrtype; /* Host address type. */
    int h_length; /* Length of address. */
    char **h_addr_list; /* List of addresses from name
                        server.*/
#define h_addr h_addr_list[0] /* The first address in
                        the address list. */
};
```

# Domain Name System (3): Host Entry Structure



# Domain Name System (4): Example of Host Entry Structure

struct hostent{



You can try the command: `nslookup v.qq.com`

# Domain Name System (5)

- Functions for retrieving host entries from DNS

- **gethostbyname**: query key is a DNS domain name.

```
#include <netdb.h>
```

```
struct hostent * gethostbyname (const char *hostname);
```

- **gethostbyaddr**: query key is an IP address.

```
#include <netdb.h>
```

```
struct hostent * gethostbyaddr (const char *addr, int len, int family);
```

↓  
**AF\_INET**  
for IPv4



---

# Socket Interface



# Sockets Interface

---

- Functions
- Definitions
- Types





# Sockets Interface – functions

---

- Created in the early 80's as part of the original Berkeley distribution of Unix that contained an early version of the Internet protocols
- Provides a user-level interface to the network
- Underlying basis for all Internet applications
- Based on **client/server** programming model



# Sockets Interface – definitions(1)

---

- What is a socket?
  - To the kernel, a socket is an **endpoint** of communication.
  - To an application, a socket is a **file descriptor** that lets the application read/write from/to the network.
    - **Remember**: All Unix I/O devices, including networks, are modeled as files.
- Clients and servers communicate with each other by **reading from and writing to socket descriptors**.
- The main distinction between regular file I/O and socket I/O is how the application “opens” the socket descriptors.



# File Model in Unix/Linux

- In Unix/Linux, all I/O devices are treated as files
  - Identified with File Descriptors
  - File operations
    - open
    - close
    - lseek
    - read
    - write
    - ...

Sample File Descriptor Table  
(One per Process)

0	stdin
1	stdout
2	stderr
3	file
4	device
5	socket
..	...

# Sockets Interface – definitions(2)

## ■ **Internet-specific** socket address (bits/socket.h)

```
struct sockaddr_in {  
    unsigned short sin_family; /* address family (always AF_INET) */  
    unsigned short sin_port; /* port num in network byte order */  
    struct in_addr sin_addr; /* IP addr in network byte order */  
    unsigned char sin_zero[8]; /* pad to sizeof(struct sockaddr) */  
};
```

一定转换成network by order

## ■ **Address family**: Domains refer to the area where the communicating processes exist. Commonly used domains include:

- AF\_UNIX: for communication between processes on one system;
- **AF\_INET** (IPv4): for communication between processes on the same or different systems using the DARPA standard protocols (IP/UDP/TCP)
- AF\_INET6 (IPv6)
- AF\_LOCAL (Unix domain)
- AF\_UNSPEC (the importance will be explained later)
- ...

# Sockets Interface – definitions(3)

- **Generic socket address** (<sys/socket.h>)

```
struct sockaddr {  
    unsigned short sa_family; /* protocol family */  
    char sa_data[14]; /* protocol-specific address,  
                        up to 14 bytes. */  
};
```

- **Protocol family**

- PF\_LOCAL: Local to host, pipes and file-domain
- PF\_UNIX: Old BSD name for PF\_LOCAL
- **PF\_INET: IP protocol family**
- PF\_AX25: Amateur radio AX.25
- PF\_IPX: Novell internet protocol
- PF\_INET6: IP version 6
- PF\_ATMSVC: ATM SVCs
- PF\_APPLETALK: Appletalk DDP
- ...



# Sockets Interface – definitions(4)


- Generic socket address **and** Internet-specific socket address
  - Pointer to generic socket address is used for address arguments to **connect()**, **bind()** and **accept()**
  - Must **cast** Internet-specific socket address (`sockaddr_in *`) to generic socket address (`sockaddr *`) for connect, bind, and accept

```
struct sockaddr_in serv;  
/* fill in serv{}*/  
bind (sockfd, (struct sockaddr *)&serv, sizeof(serv));
```

# Socket Address

**Generic socket address**

```
struct sockaddr {  
    unsigned short sa_family;    /* PF_INET for IPv4 */  
    char sa_data[14];           /* protocol-specific address,  
                                up to 14 bytes. */  
};
```



**Internet-specific socket address**

```
struct sockaddr_in {  
    unsigned short sin_family;    /* AF_INET */  
    unsigned short sin_port;      /* 16-bit port number */  
    struct in_addr sin_addr;      /* Network Byte Order */  
    char sin_zero[8];            /* 32-bit IP Address */  
    /* Network Byte Order */  
};
```



# Sockets Interface – types(1)

---

## ■ Stream Socket

- Service: **reliable** (i.e. sequenced, non-duplicated, non-corrupted) bidirectional delivery of byte-stream data
- Metaphor: a phone call
- `int s = socket (PF_INET, SOCK_STREAM, 0);`

## ■ Datagram Socket

- Service: **unreliable**, unsequenced datagram
- Metaphor: sending a letter
- `int s = socket (PF_INET, SOCK_DGRAM, 0);`

## ■ Raw Sockets Service

- allows user-defined protocols that interface with IP
- Requires `root` access
- Metaphor: playing with an erector set
- `int s = socket (PF_INET, SOCK_RAW, protocol);`

- **SOCK\_STREAM and SOCK\_DGRAM are the most common types of sockets used within UNIX/Linux**





# Sockets Interface – types(2)

---

- Reliably-delivered Message Socket
  - Service: reliable datagram
  - Metaphor: sending a registered letter
  - Similar to datagram socket but ensure the arrival of the datagrams
  - `int s = socket (PF_NS, SOCK_RDM, 0);`
- Sequenced Packet Stream Socket
  - Service: reliable, bi-directional delivery of recordoriented data
  - Metaphor: record-oriented TCP
  - Similar to stream socket but using fixed-size datagrams
  - `int s = socket (PF_NS, SOCK_SEQPACKET, 0);`



---

# Major System Calls



# Socket Programming: Telephone Analogy

---

- A telephone call over a “telephony network” works as follows:
  - Both parties have a telephone installed.
  - A phone number is assigned to each telephone.
  - Turn on ringer to listen for a caller.
  - Caller lifts telephone and dials a number.
  - Telephone rings and the receiver of the call picks it up.
  - Both Parties talk and exchange data.
  - After conversation is over they hang up the phone.



# Dissecting the Analogy

---

- A network application works as follows:
  - An **endpoint** (telephone) for communication is created on both ends.
  - An **address** (phone no) is assigned to both ends to distinguish them from the rest of the network.
  - One of the endpoint (receiver) **waits** for the communication to start.
  - The other endpoint (caller) **initiates** a connection.
  - Once the call has been **accepted**, a connection is made and data is exchanged (talk).
  - Once data has been exchanged the **endpoints are closed** (hang up).



# In the world of sockets.....

---

- `socket()` – Create endpoint for communication
- `bind()` - Assign a unique telephone number
- `listen()` – Wait for a caller
- `connect()` - Dial a number
- `accept()` – Receive a call
- `send()`, `recv()` – Talk
- `close()` – Hang up



# System Calls

---

- Socket operation
- Byte order operation
- Address formats conversion
- Socket option
- Name and address operation



# System Calls – Socket Operation

---

- **socket()**
  - returns a socket descriptor
- **bind()**
  - What address I am on / what port to attach to
- **connect()**
  - Connect to a remote host
- **listen()**
  - Waiting for someone to connect to my port
- **accept()**
  - Get a socket descriptor for an incoming connection
- **send() and recv()**
  - Send and receive data over a connection
- **read(), write()**
  - Read from / Write to a particular socket, similar to recv()/ send()
- **sendto() and recvfrom()**
  - Send and receive data without connection
- **close() and shutdown()**
  - Close a connection Two way / One way



# System Calls – Byte Order Conversion

---

- `htonl()`
  - host byte order → network byte order for *long int*
- `htons()`
  - host byte order → network byte order for *short int*
- `ntohl()`
  - network byte order → host byte order for *long int*
- `ntohs()`
  - network byte order → host byte order for *short int*





# System Calls – Address Formats Conversion

---

- **inet\_aton()**
  - IP address in **numbers-and-dots notation** (ASCII string) → IP address structure in **network byte order**
- **inet\_addr()**
  - same function with inet\_aton()
- **inet\_ntoa()**
  - IP address structure in **network byte order** → IP address in **numbers-and-dots notation** (ASCII string)
- **inet\_pton()**
  - Similar to inet\_aton() but working with **IPv4 and IPv6**
- **inet\_ntop()**
  - Similar to inet\_ntoa() but working with **IPv4 and IPv6**



# System Calls – Socket Option

---

- `getsockopt()`
  - Allow an application to require information about the socket
- `setsockopt()`
  - Allow an application to set a socket option
- eg. get/set sending/receiving buffer size of a socket



## System Calls – Name and Address Operation

---

### **gethostbyname()**

- retrieving host entries from DNS and the query key is a DNS domain name

### **gethostbyaddr()**

- retrieving host entries from DNS and the query key is an IP address

### **gethostname()**

- Obtaining the name of a host

### **getservbyname()**

- Mapping a named service onto a port number

### **getservbyport()**

- Obtaining an entry from the services database given the port number assigned to it

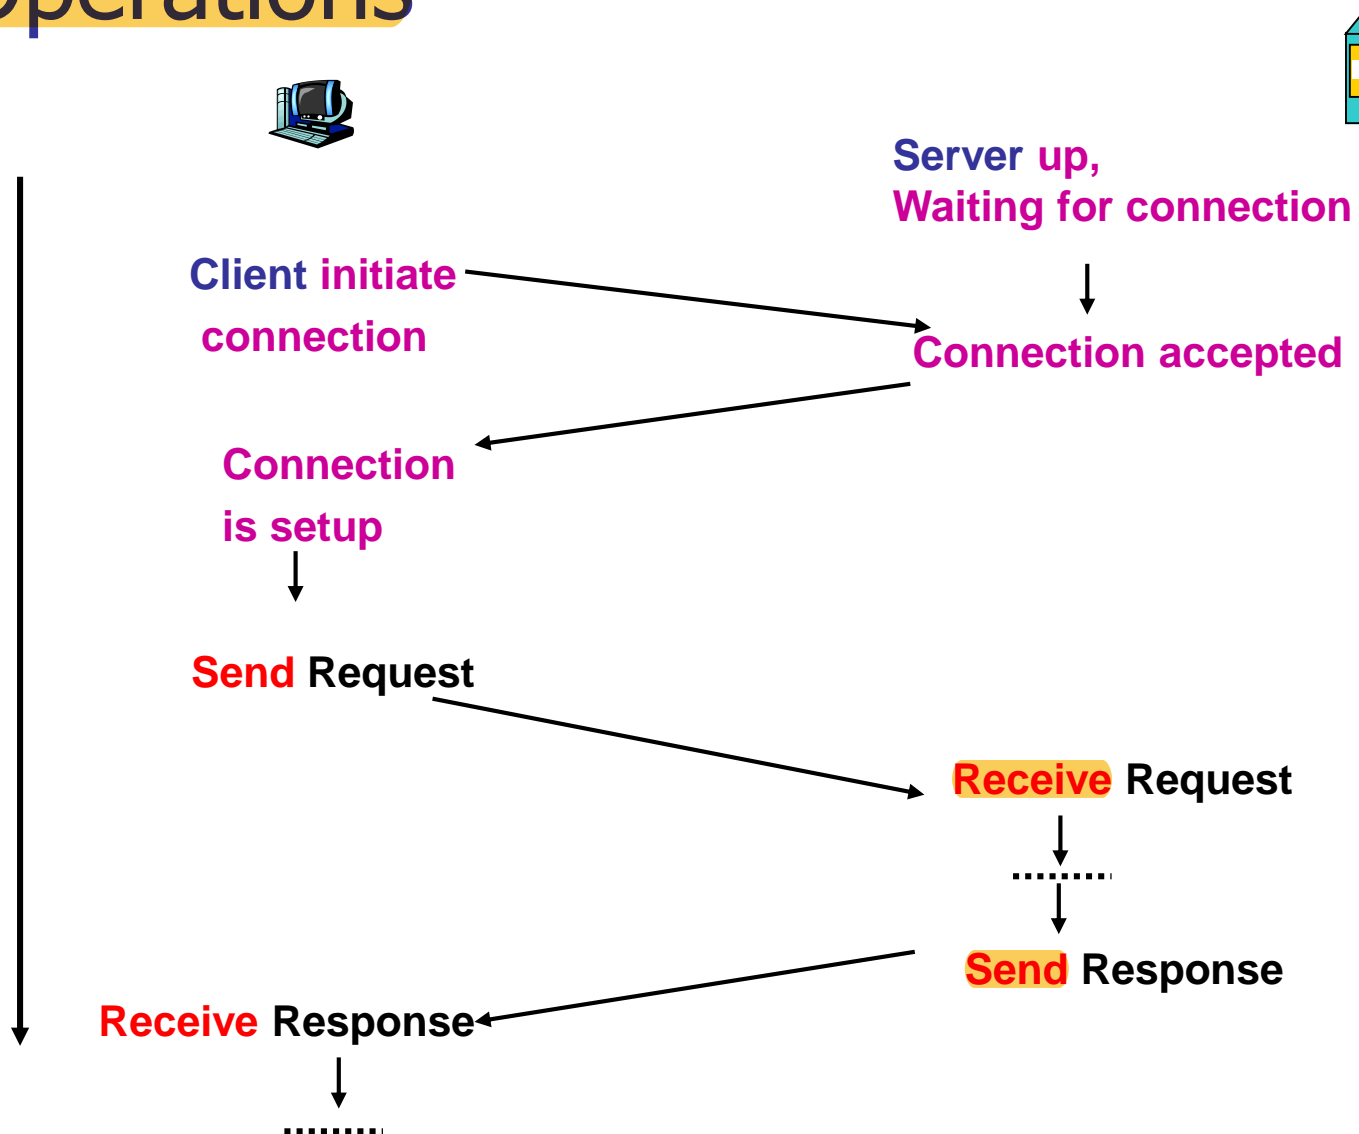


***Using UDP to query Local DNS Server***

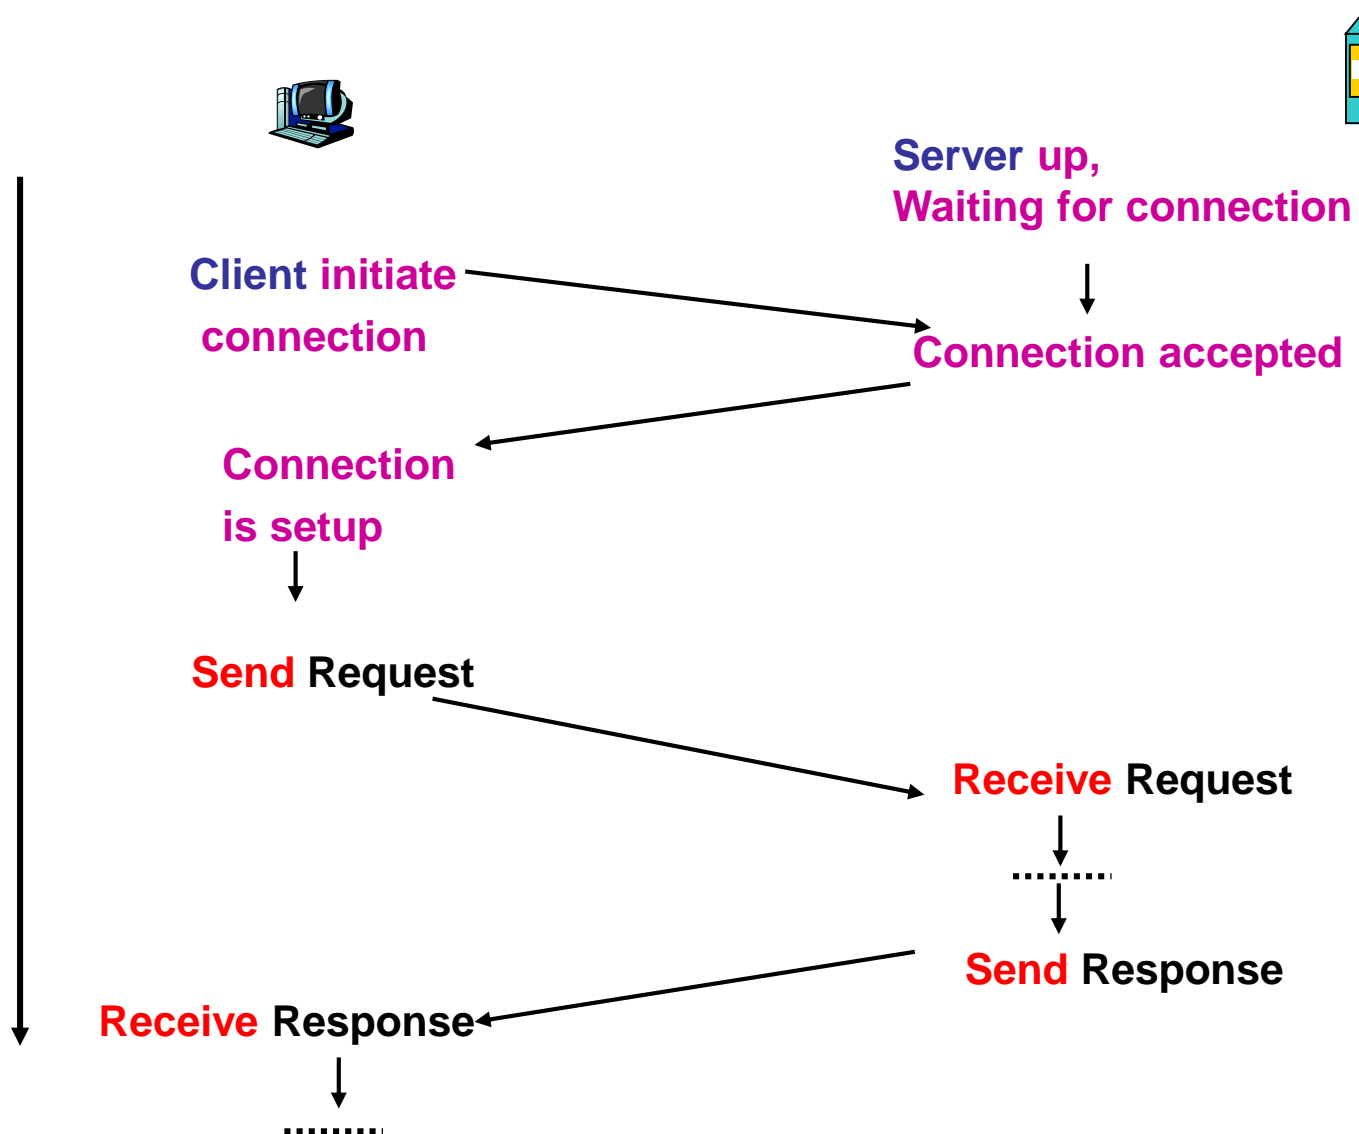
***Returns a pointer to struct hostent (host entry structure) on success***

# Process of Socket Operation:

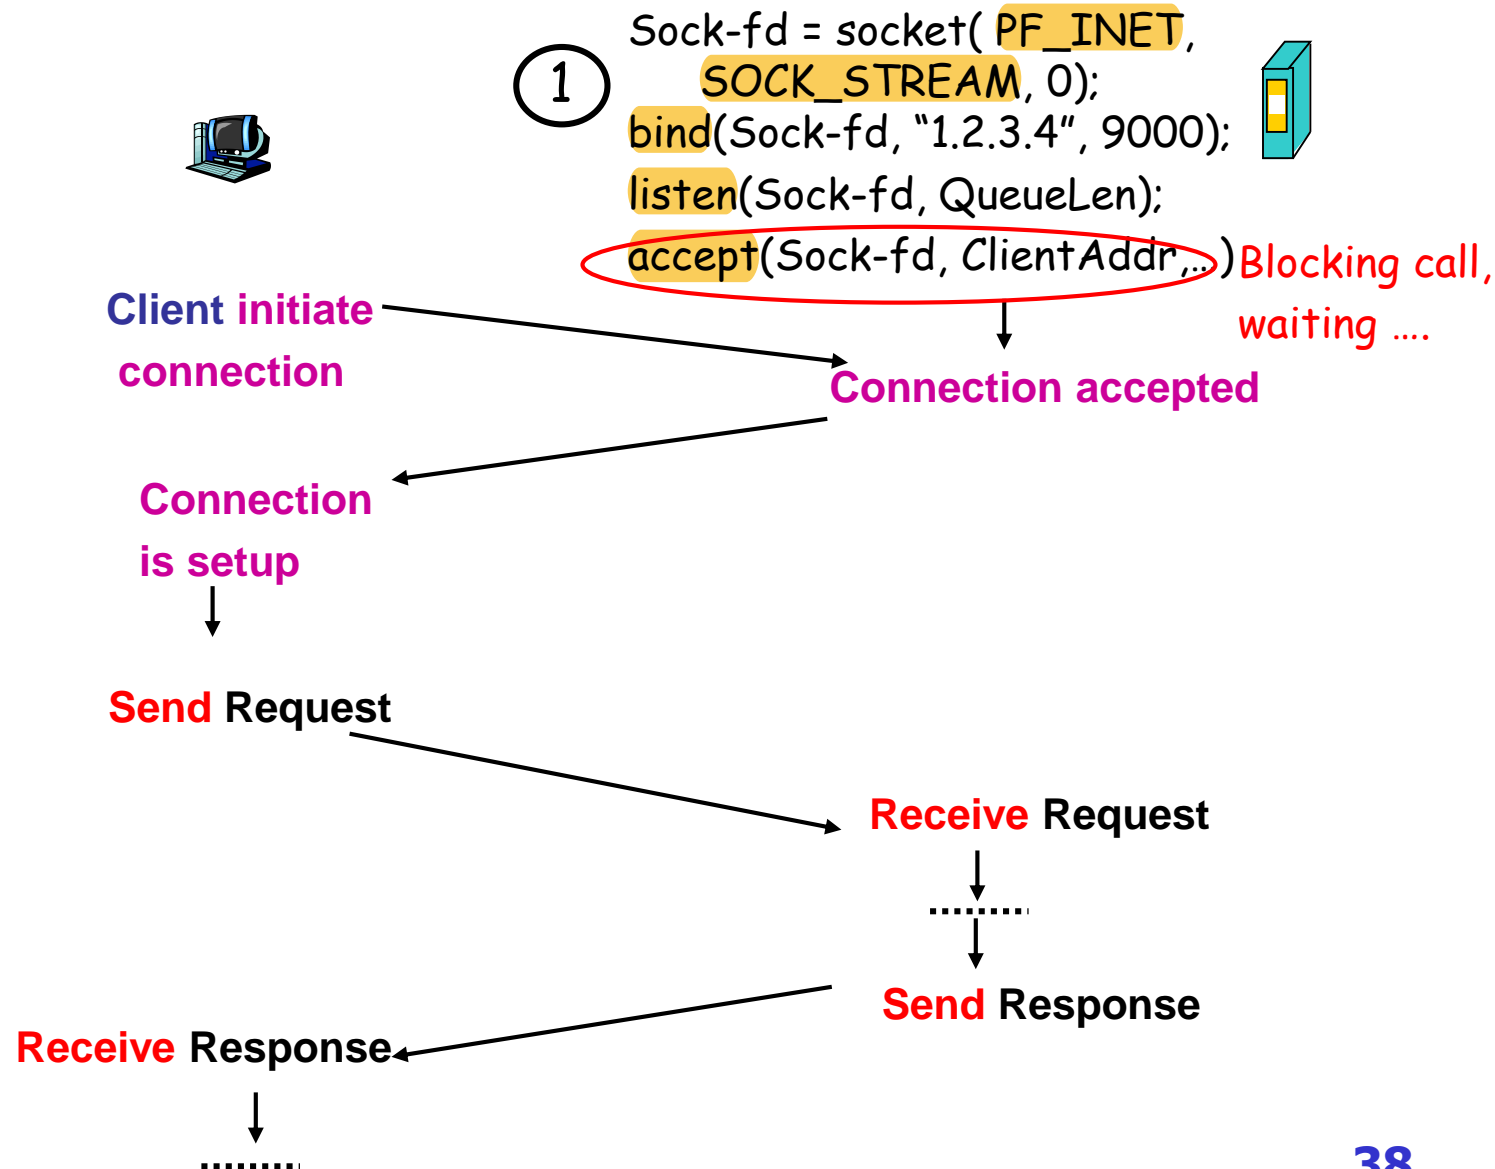
## TCP Operations



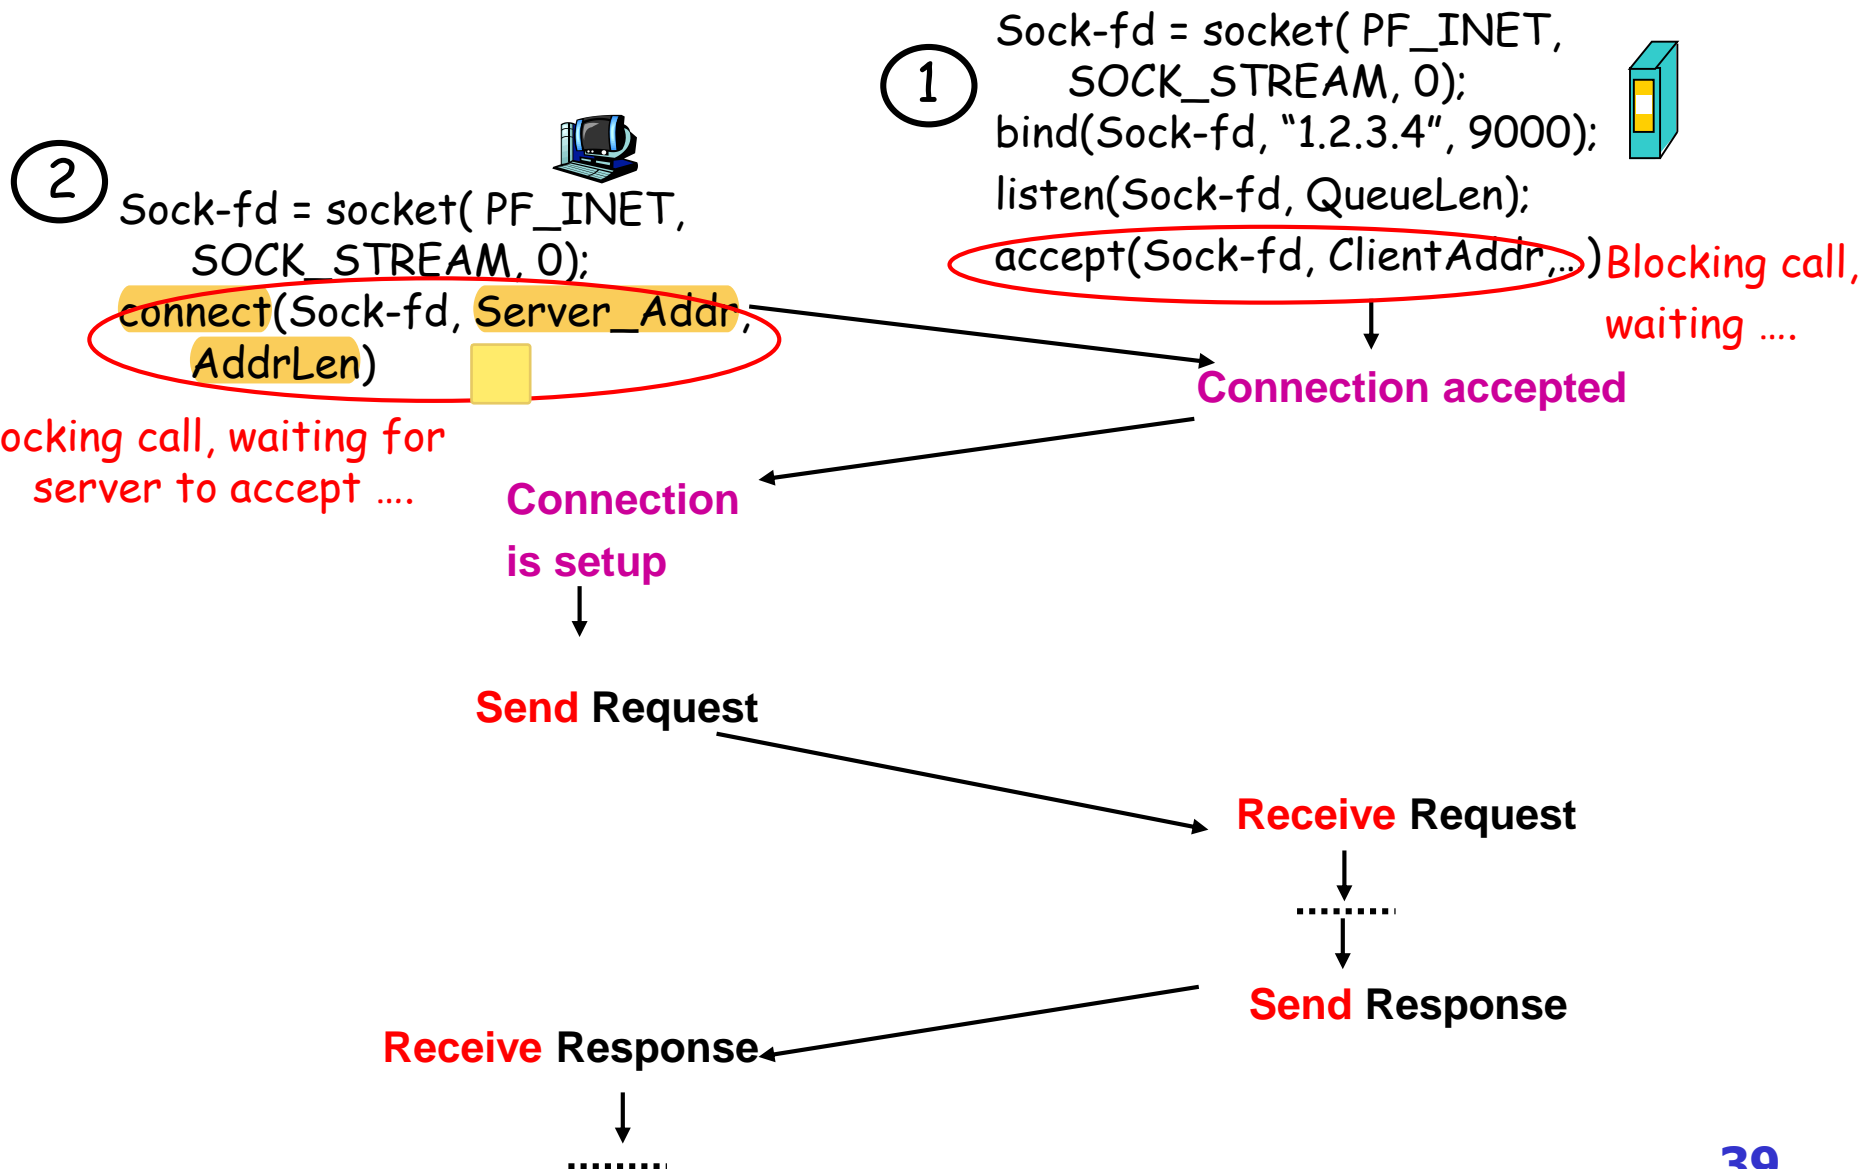
# Map to TCP socket programming



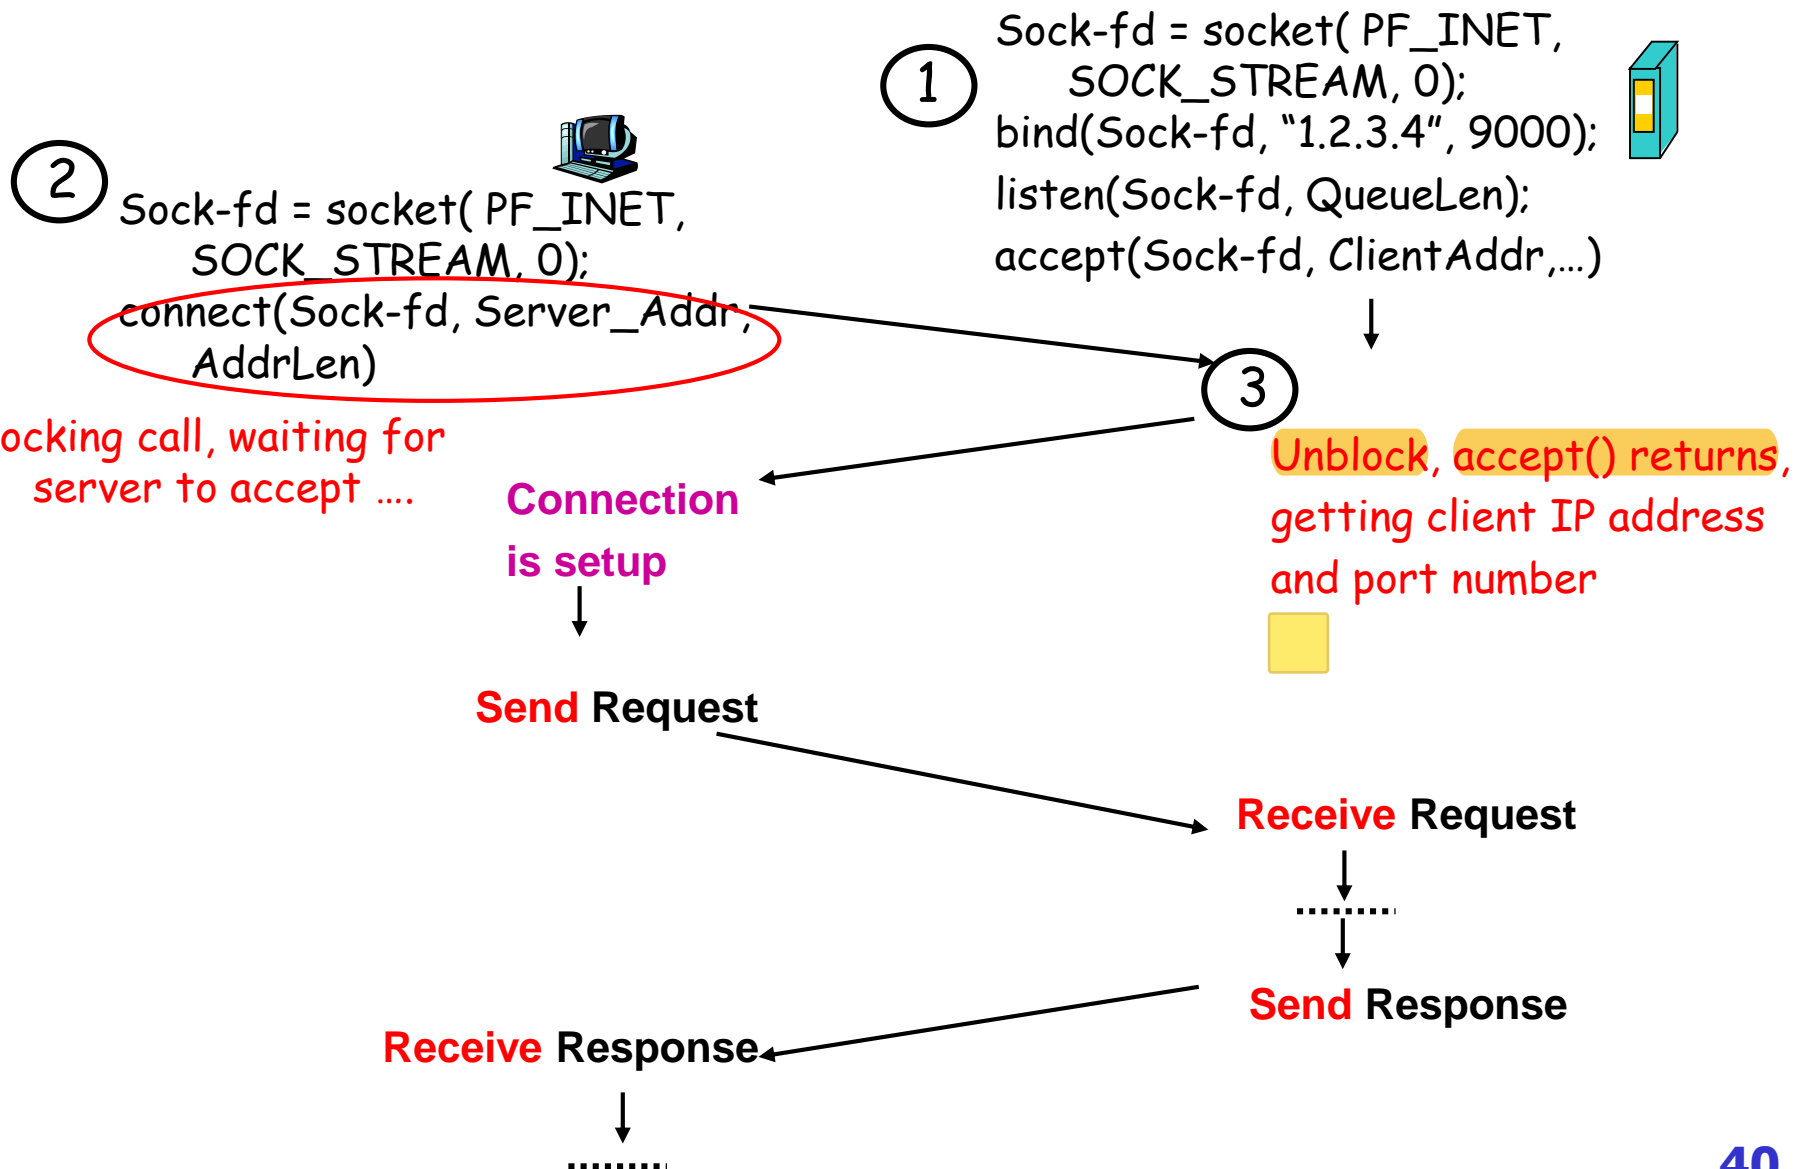
# Map to TCP socket programming



# Map to TCP socket programming

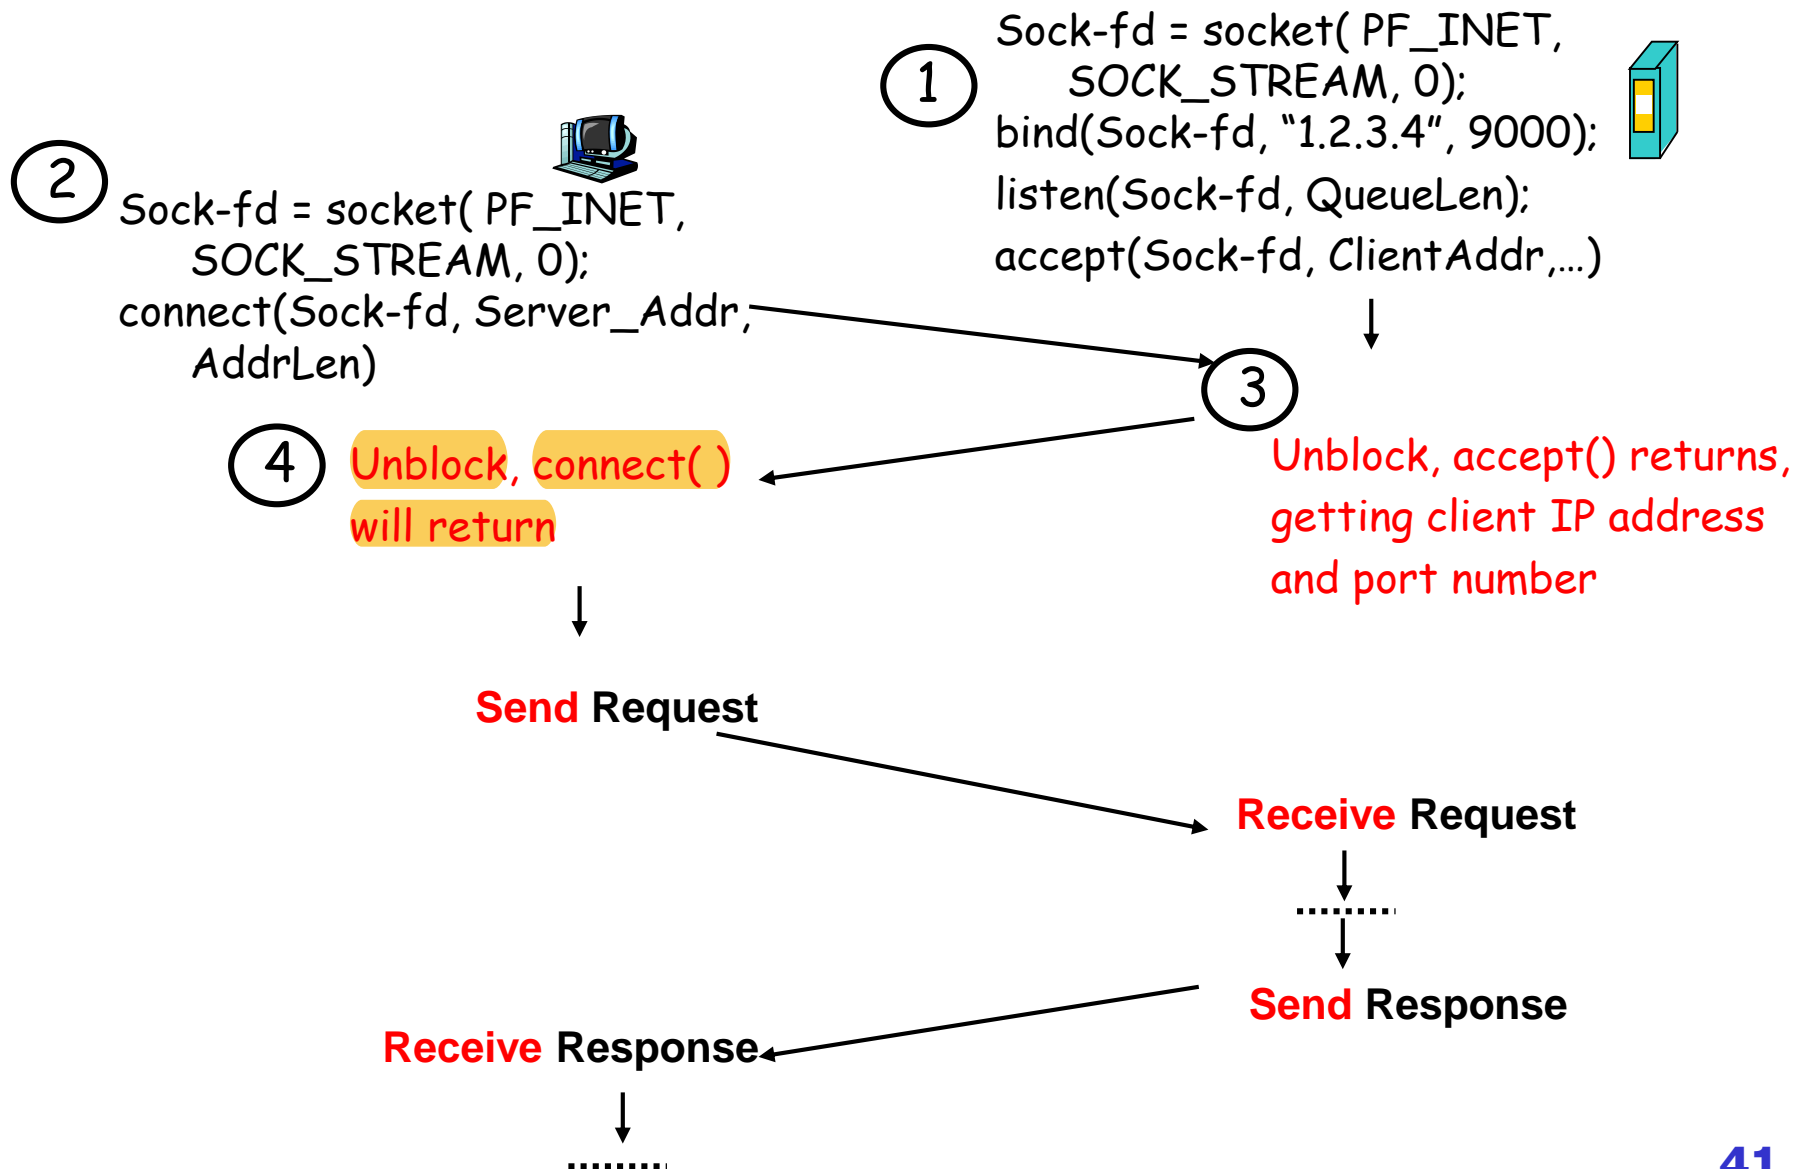


# Map to TCP socket programming

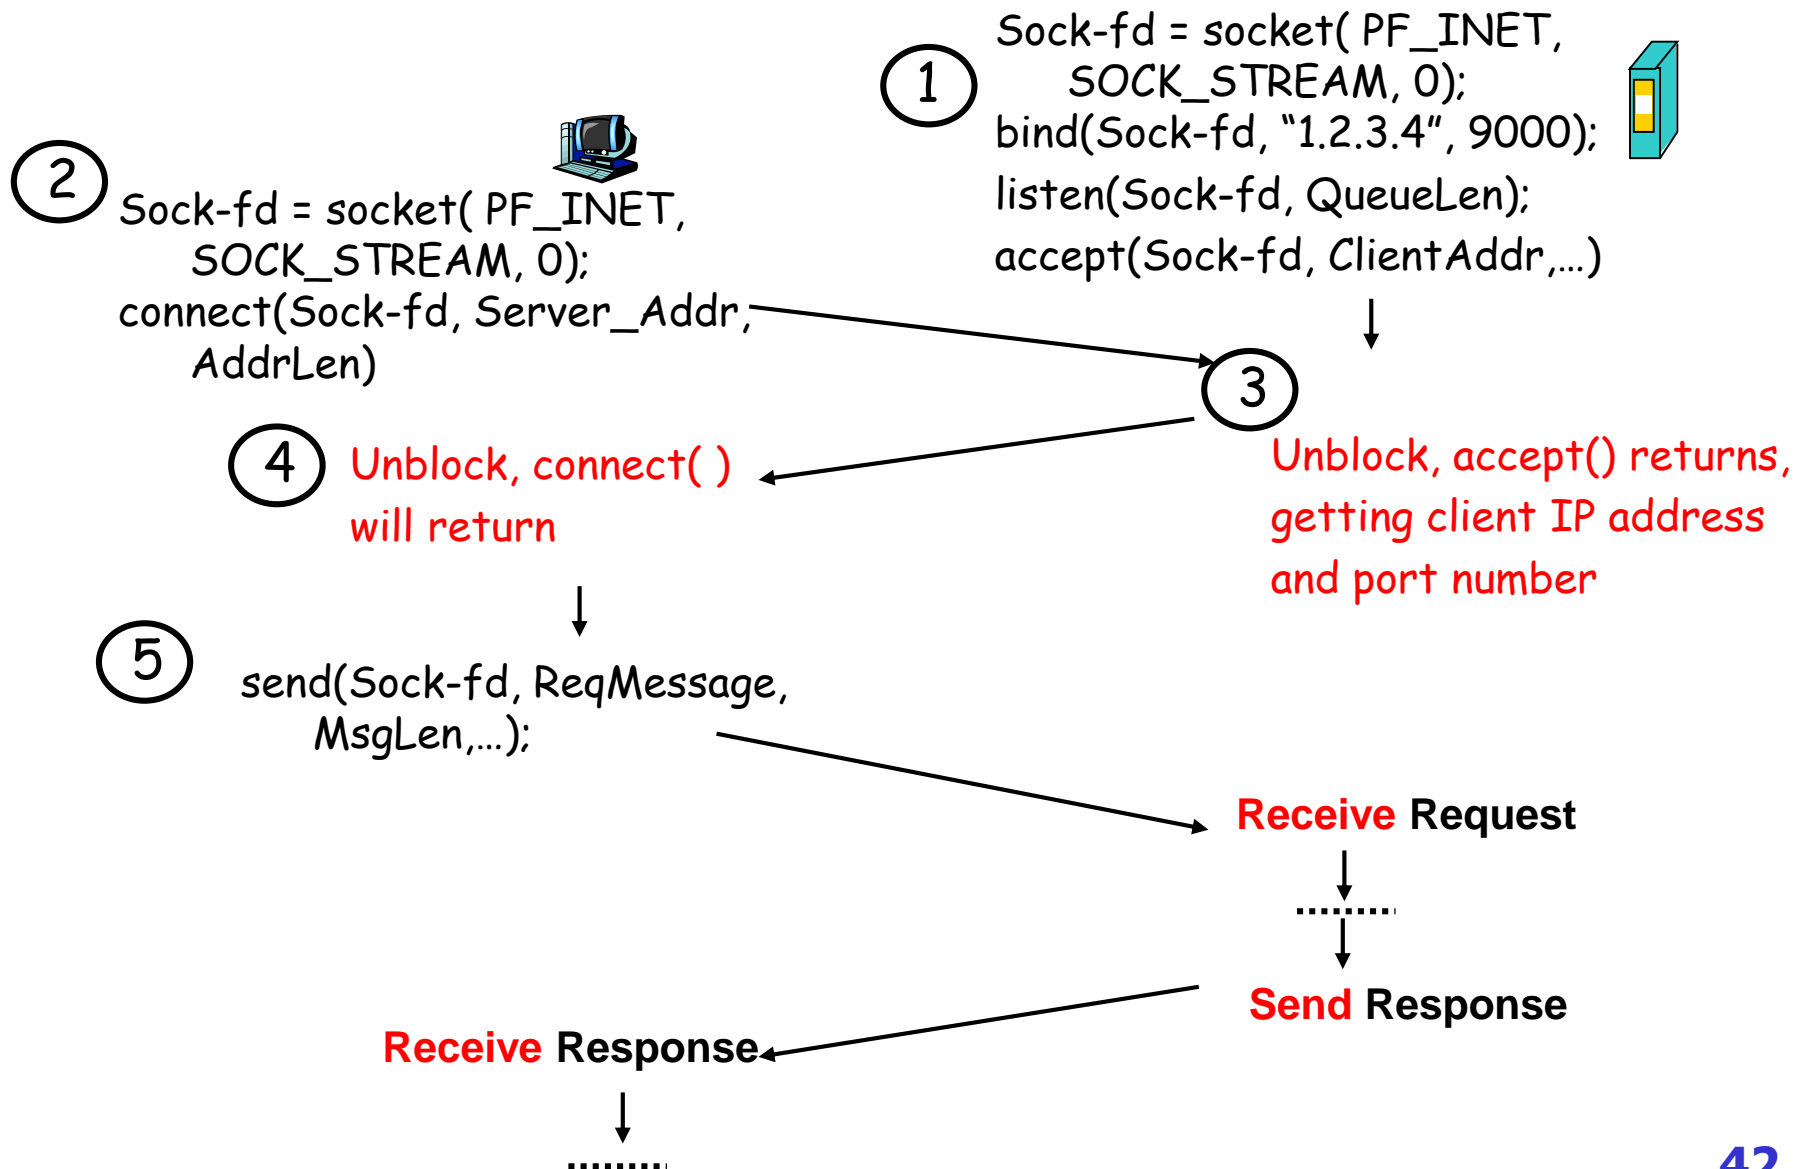




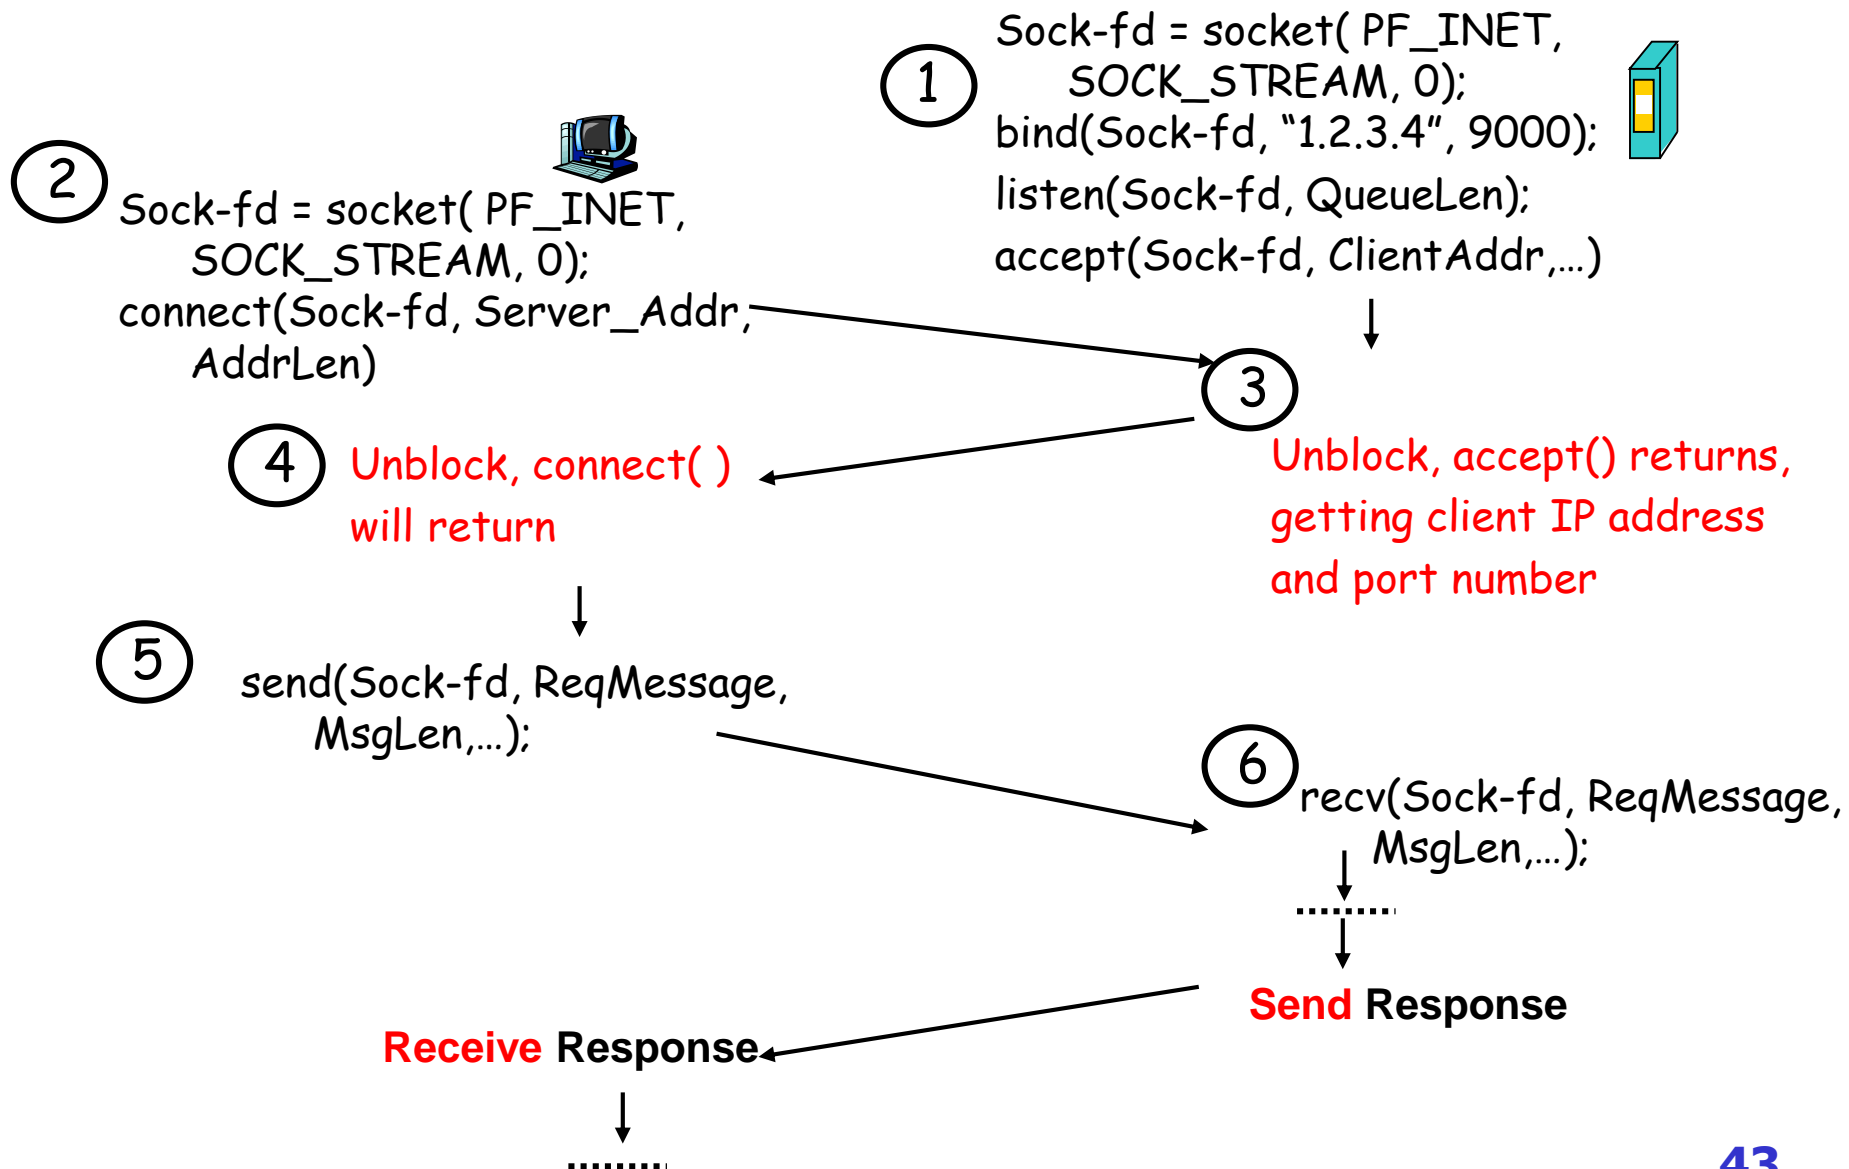
# Map to TCP socket programming



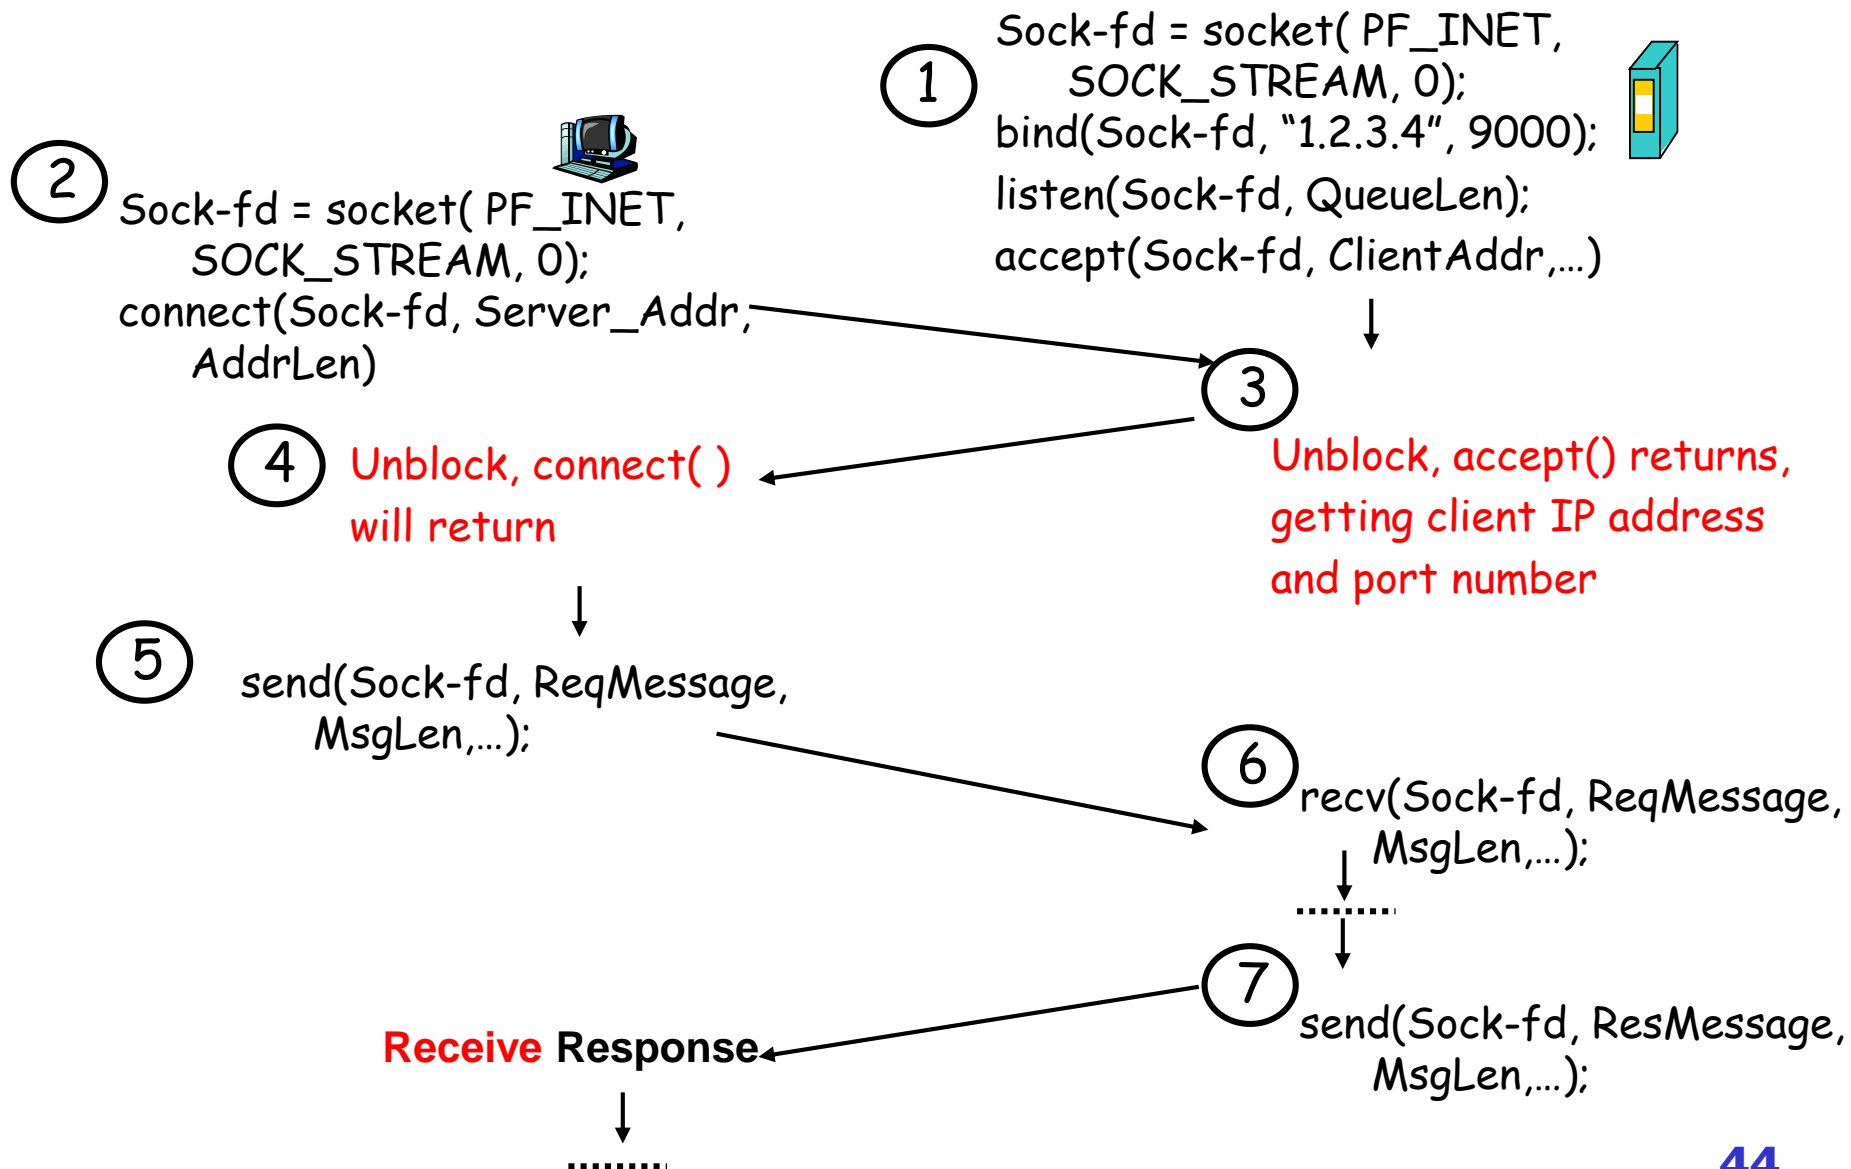
# Map to TCP socket programming



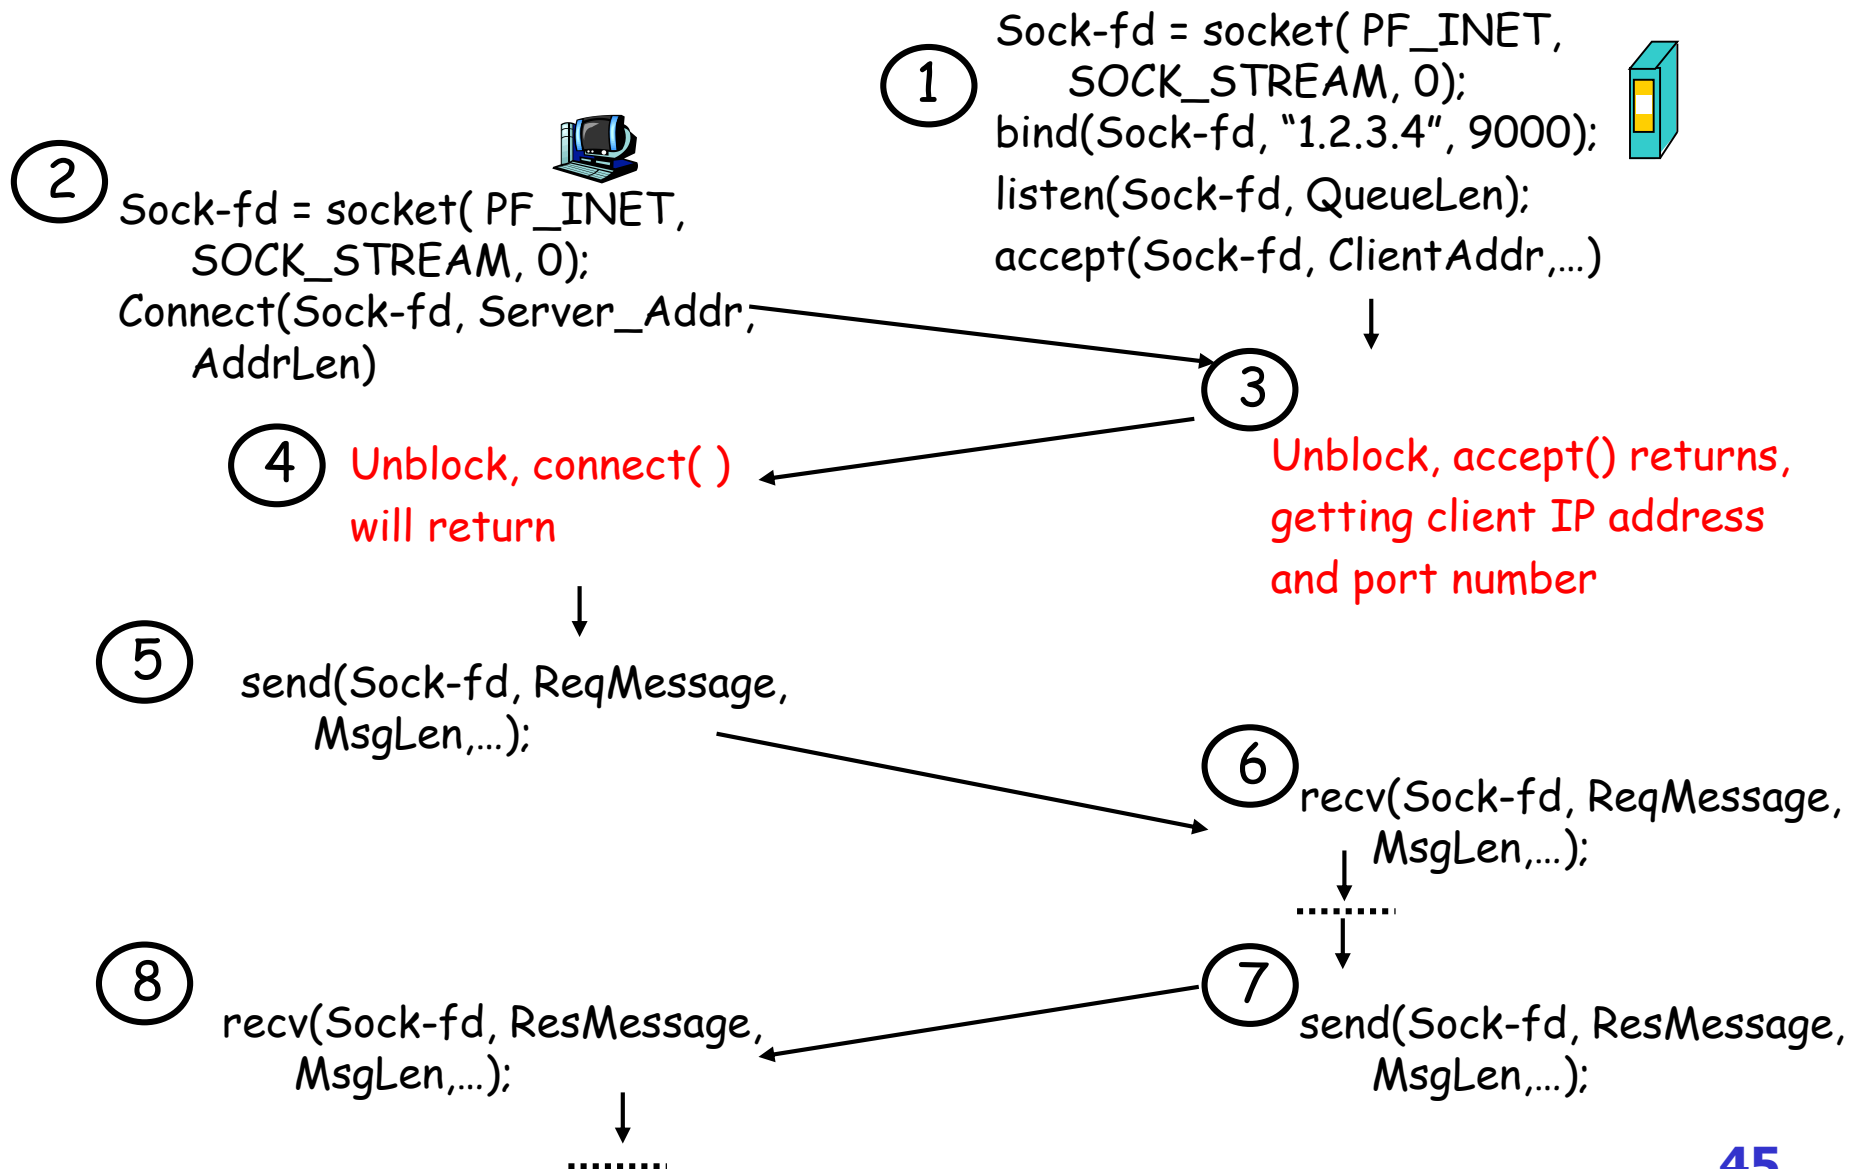
# Map to TCP socket programming



# Map to TCP socket programming



# Map to TCP socket programming



# System Calls – socket()

- An application calls socket() to create a new socket that can be used for network communication
- The call returns a descriptor for the newly created socket

```
#include <sys/socket.h>
```

```
int socket (int family, int type, int protocol);
```

- **Socket descriptor**
  - success
  - -1 – failed

## Protocol family

**PF\_INET**  
PF\_INET6  
PF\_UNIX  
.....

## Socket type

**SOCK\_STREAM**  
**SOCK\_DGRAM**  
SOCK\_RAW  
.....

## Specific protocol

Often be set as **0**  
Be not 0 in raw socket

# System Calls – bind()

- An application calls `bind()` to specify the local endpoint address (a **local IP address and protocol port number**) for a socket
- For TCP/IP, the endpoint address uses the ***sockaddr\_in*** structure.
- Must **cast** Internet-specific socket address (`struct sockaddr_in *`) to generic socket address (`struct sockaddr *`) for **bind**
- **Servers** use ***bind()*** to **specify** the **well-known port** at which they will await connections

```
#include <sys/socket.h>
```

```
int bind (int sockfd, const struct sockaddr *myaddr, socklen_t addrlen);
```

0 – success  
-1 – failed

Socket  
descriptor to  
be bound

The local address to  
which the socket should  
be bound

The length of the  
structure (bytes)

# System Calls – listen()

- Connection-oriented **servers** call *listen()* to place a socket in *passive mode* and make it ready to accept incoming connections
- *listen()* also sets the number of incoming connection requests that the protocol software should enqueue for a given socket while the server handles another request
- It only applies to socket used with **TCP**

```
#include <sys/socket.h>
```

```
int listen (int sockfd, int qlength);
```

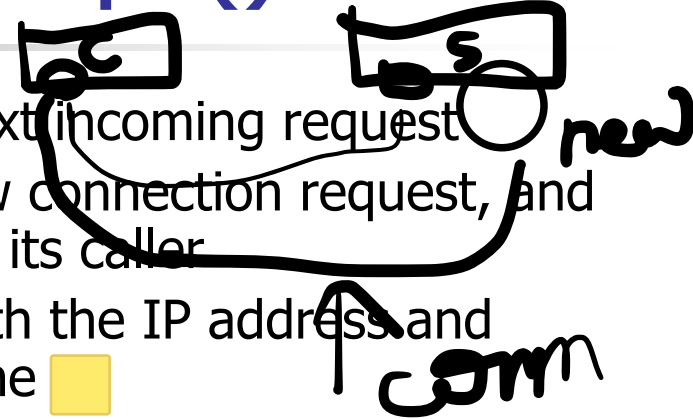
0 – success  
-1 – failed

Socket that should  
be prepared for  
use by a server

The length of the  
request queue for  
that socket



# System Calls – accept()

- The **server** calls `accept()` to extract the next incoming request
  - `accept()` **creates a new socket** for each new connection request, and returns the descriptor of the new socket to its caller
  - `accept()` fills in the structure (`sockaddr`) with the IP address and protocol port number of the remote machine
  - **Must cast** Internet-specific socket address (`struct sockaddr_in *`) to **generic** socket address (`struct sockaddr *`) for `accept()`
- 

```
#include <sys/socket.h>
```

```
int accept (int sockfd, struct sockaddr *cliaddr, socklen_t *addrlen);
```

Socket Descriptor  
(non-zero) –  
success  
-1 – failed

Socket on  
which to  
wait

The address of the client  
that placed the request

The length of the  
client address

new socket descriptor

# System Calls – connect()

- After creating a socket, a **client** calls *connect()* to establish an active connection to a remote server
- Must **cast** Internet-specific socket address (`struct sockaddr_in *`) to generic socket address (`struct sockaddr *`) for **connect**

```
#include <sys/socket.h>
```

```
int connect (int sockfd, const struct sockaddr *servaddr, socklen_t addrlen );
```

0 – success  
-1 – failed

**Socket** to  
**connect**

The **server's** address to  
which the socket should  
connect

The **length** of the  
server's address  
(bytes)

# System Calls – send()

- Both clients (to transmit request) and servers (to transmit replies) used *send()* to transfer data across a **TCP connection**
- The application passes the descriptor of a socket to which the data should be sent, the address of the data to be sent, and the length of the data
- Usually, send copies outgoing data into buffers in the OS kernel

```
#include <sys/socket.h>
```

```
ssize_t send (int sockfd, const void *buff, size_t nbytes, int flags);
```

Number of bytes  
that are sent  
successfully

-1 – failed

socket  
descriptor  
you want to  
send data  
to

The address  
of the data  
to be sent

The number  
of bytes to  
be sent

The flag controlling  
the connection,  
usually 0

# System Calls – `recv()`

- Both clients (to receive a reply) and servers (to receive a request) use *recv* to receive data from a **TCP connection**
- If the buffer cannot hold an incoming user datagram, *recv* fills the buffer and discards the remainder

```
#include <sys/socket.h>
```

```
ssize_t recv (int sockfd, void *buff, size_t nbytes, int flags);
```

Number of bytes  
that are received  
successfully  
-1 – failed

socket  
descriptor  
to read  
from

The address in  
memory into  
which the data  
to be placed

The  
length of  
the buffer  
area

The flag  
controlling the  
reception,  
usually 0

# System Calls – `sendto()` & `recvfrom()`

- Allow the caller to send or receive a message over **UDP**
- `sendto()` requires the caller to specify a destination
- `recvfrom()` uses an argument to specify where to record the sender's address

```
#include <sys/socket.h>
ssize_t sendto (int sockfd, const void *buff, size_t nbytes, int flags,
               const struct sockaddr *to, socklen_t addrlen);
```

Destination  
address

The length of the  
destination address

```
#include <sys/socket.h>
ssize_t recvfrom (int sockfd, void *buff, size_t nbytes, int flags,
                 struct sockaddr *from, socklen_t *addrlen);
```

Number of bytes that are sent  
or received successfully  
-1 – failed

Where to record the  
sender's address

The length of the  
sender's address



# Using Read and Write with sockets

---

- In Linux, as in most other UNIX systems, programmers can use *read* instead of *recv*, and *write* instead of *send*
  - *int read (sockfd, bptr, buflen)*
  - *int write (sockfd, bptr, buflen)*
- The chief advantage of *send* and *recv* is that they are easier to spot in the code



# System Calls – close()

---

- Once a client or server finishes using a socket, it calls *close* to deallocate it
- Any unread data waiting at the socket will be discarded

```
#include <unistd.h>  
int close (int sockfd);
```

0 – success  
-1 – failed

Socket to be closed

# System Calls – inet\_aton() & inet\_addr()

- Converts an IP address in numbers-and-dots notation into unsigned long in **network byte order**

```
#include <arpa/inet.h>
```

```
int inet_aton (const char *string, struct in_addr *address);
```

1 – success  
0 – error

Pointer to the string that  
contains the numbers-  
and-dots notation

Pointer to IP address  
structure

```
#include <arpa/inet.h>
```

```
in_addr_t inet_addr (const char *string);
```

When success: return the 32-bit  
address in network byte order  
When failed: return INADDR\_NONE

Pointer to the string that  
contains the numbers-and-  
dots notation






# System Calls – inet\_ntoa()

- Mapping a 32-bit integer (an IP address in **network byte order**) to an ASCII string in dotted decimal format

```
#include <arpa/inet.h>  
char *inet_ntoa (struct in_addr inaddr);
```

Pointer to the string  
contains the numbers-  
and-dots notation


IP address structure  
in network byte order




# Sample Programs

# Overview of TCP-based sockets API

## TCP Client



## TCP Server



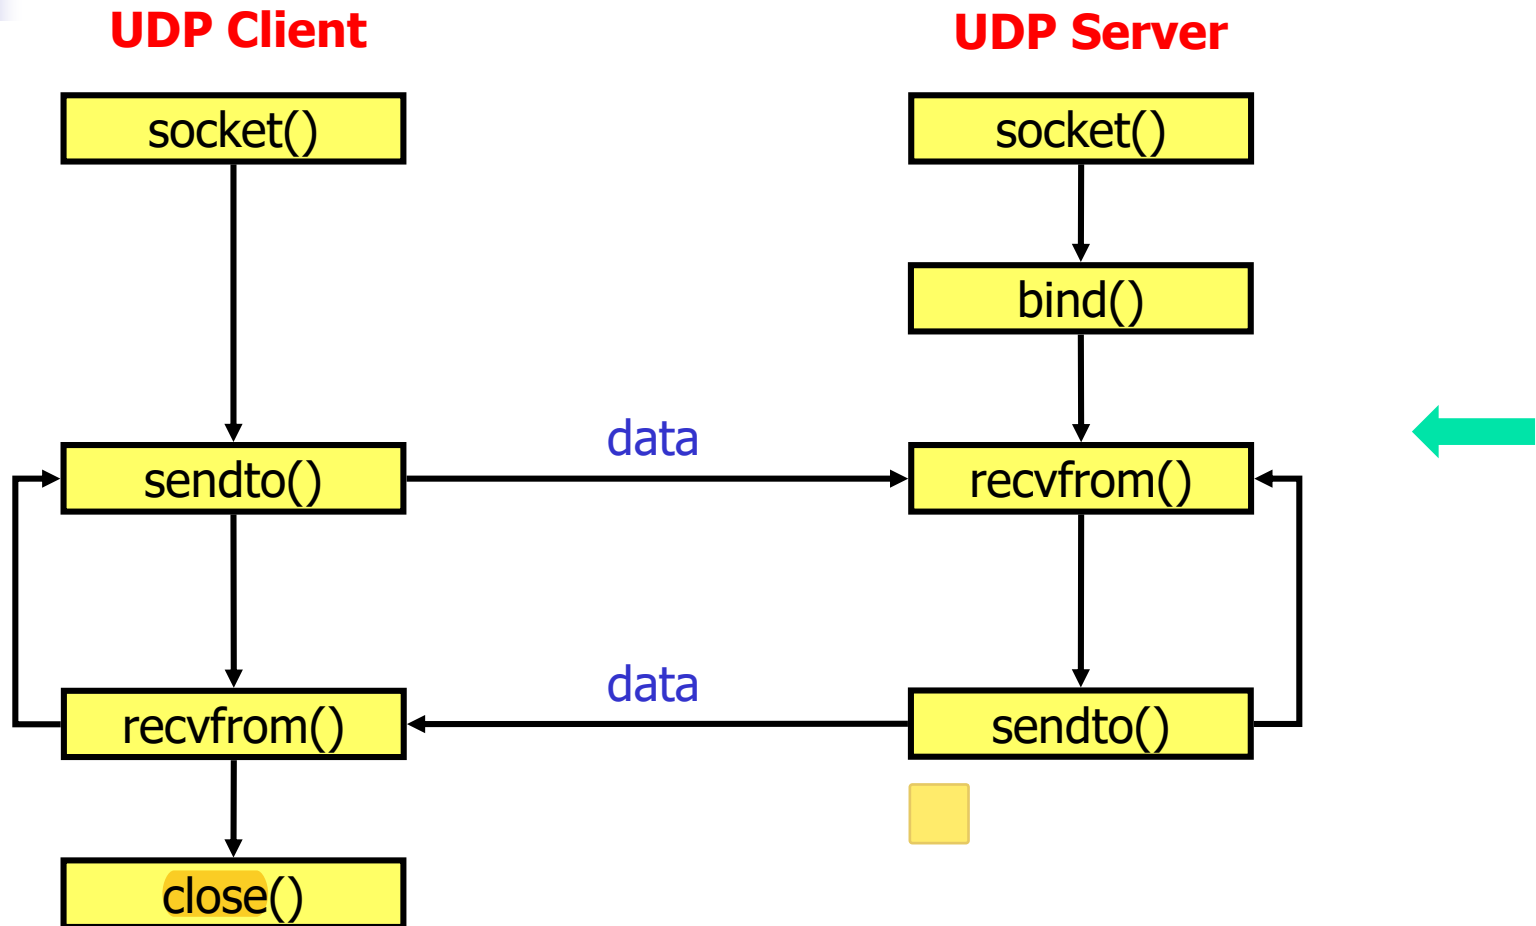
Connection request

data

data

Await  
connection  
request from  
next client

# Overview of UDP-based sockets API





# Sample programs


---

- UDP-based echo service
  - An echo service simply sends back to the originating source any data it receives
  - A very useful debugging and measurement tool
  - UDP Based Echo Service: be defined as a datagram based application on UDP. A server listens for UDP datagrams on UDP port 7. When a datagram is received, the data from it is sent back in an answering datagram.
- Sample programs
  - `udpechoclt.c`
  - `udpechosvr.c`



# Basic flow of UDP-based echo service

---





# Head part of UDP EchoClient

---

```
#include <stdio.h> /* for printf() and fprintf() */
#include <sys/socket.h> /* for socket(), sendto() and
                        recvfrom() */
#include <arpa/inet.h> /* for sockaddr_in and inet_addr() */
#include <stdlib.h> /* for atoi() and exit() */
#include <string.h> /* for memset() */
#include <unistd.h> /* for close() */
```



# Initial part of UDP EchoClient

---

```
#define ECHOMAX 255 /* Longest string to echo */

int main(int argc, char *argv[])
{
    int sock; /* Socket descriptor */
    struct sockaddr_in echoServAddr; /* Echo server address */
    struct sockaddr_in fromAddr; /* Source address of echo */
    unsigned short echoServPort; /* Echo server port */
    unsigned int fromSize; /* In-out of address size
                           for recvfrom() */
    char *servIP; /* IP address of server */
    char *echoString; /* String to send to echo server */
    char echoBuffer[ECHOMAX+1]; /* Buffer for receiving
                                echoed string */
    int echoStringLength; /* Length of string to echo */
    int respStringLength; /* Length of received response */
}
```



# Argument check part of UDP EchoClient

```
if ((argc < 3) || (argc > 4)) /* Test for correct number of
arguments */
{
    printf("Usage: %s <Server IP> <Echo Word> [<Echo Port>]\n",
           argv[0]);
    exit(1);
}

servIP = argv[1]; /* First arg: server IP address (dotted quad) */
echoString = argv[2]; /* Second arg: string to echo */
if ((echoStringLen = strlen(echoString)) > ECHOMAX) /* Check input
length */
    printf("Echo word too long.\n");
if (argc == 4)
    echoServPort = atoi(argv[3]); /* Use given port, if any */
else
    echoServPort = 7; /* 7 is the well-known port for echo service */
```

ASCII to integer

# I/O part of UDP EchoClient

```
/* Create a datagram/UDP socket */
if ((sock = socket(PF_INET, SOCK_DGRAM, IPPROTO_UDP)) < 0)
    printf("socket() failed.\n");

/* Construct the server address structure */
memset(&echoServAddr, 0, sizeof(echoServAddr)); /*Zero out structure*/
echoServAddr.sin_family = AF_INET; /* Internet addr family */
echoServAddr.sin_addr.s_addr = inet_addr(servIP); /*Server IP address*/
echoServAddr.sin_port = htons(echoServPort); /* Server port */

/* Send the string to the server */
if ((sendto(sock, echoString, echoStringLength, 0,
    (struct sockaddr *) &echoServAddr, sizeof(echoServAddr)))
    != echoStringLength)
    printf("sendto() sent a different number of bytes than expected.\n");

/* Recv a response */
fromSize = sizeof(fromAddr);
if ((respStringLength = recvfrom(sock, echoBuffer, ECHOMAX, 0,
    (struct sockaddr *) &fromAddr, &fromSize)) != echoStringLength)
    printf("recvfrom() failed\n");
```

Generic socket address



# Last part of UDP EchoClient

---

```
if (echoServAddr.sin_addr.s_addr != fromAddr.sin_addr.s_addr)
{
    printf("Error: received a packet from unknown source.\n");
    exit(1);
}

/* null-terminate the received data */
echoBuffer[respStringLen] = '\0';
printf("Received: %s\n", echoBuffer); /*Print the echoed message*/
close(sock);
exit(0);
}
```



# Head part of UDP EchoServer

---

```
#include <stdio.h> /* for printf() and fprintf() */
#include <sys/socket.h> /* for socket(), bind(), sendto()
                        and recvfrom() */
#include <arpa/inet.h> /* for sockaddr_in and inet_ntoa() */
#include <stdlib.h> /* for atoi() and exit() */
#include <string.h> /* for memset() */
#include <unistd.h> /* for close() */
```



# Initial part of UDP EchoServer

---

```
#define ECHOMAX 255 /* Longest string to echo */

int main(int argc, char *argv[])
{
    int sock; /* Socket */
    struct sockaddr_in echoServAddr; /* Local address */
    struct sockaddr_in echoClntAddr; /* Client address */
    unsigned int cliAddrLen; /* Length of client address */
    char echoBuffer[ECHOMAX]; /* Buffer for echo string */
    unsigned short echoServPort; /* Server port */
    int recvMsgSize; /* Size of received message */
```



## Argument check part of UDP EchoServer

---

```
if (argc != 2)
{
    printf("Usage: %s <UDP SERVER PORT>\n", argv[0]);
    exit(1);
}
```



# Socket part of UDP EchoServer

```
echoServPort = atoi(argv[1]); /* First arg: local port */

/* Create socket for sending/receiving datagrams */
if ((sock = socket(PF_INET, SOCK_DGRAM, 0)) < 0)
    printf("socket() failed.\n");
/* Construct local address structure */
memset(&echoServAddr, 0, sizeof(echoServAddr));
echoServAddr.sin_family = AF_INET;
echoServAddr.sin_addr.s_addr = htonl(INADDR_ANY);
echoServAddr.sin_port = htons(echoServPort);
/* Bind to the local address */
if ((bind(sock, (struct sockaddr *) &echoServAddr,
    sizeof(echoServAddr))) < 0)
    printf("bind() failed.\n");
```



# Main loop of UDP EchoServer

---

```
for (;;) /* Run forever */
{
    /* Set the size of the in-out parameter */
    cliAddrLen = sizeof(echoClntAddr);
    /* Block until receive message from a client */
    if ((recvMsgSize = recvfrom(sock, echoBuffer, ECHOMAX,
        0, (struct sockaddr *) &echoClntAddr, &cliAddrLen)) < 0)
        printf("recvfrom() failed.\n");
    printf("Handling client %s\n", inet_ntoa(echoClntAddr.sin_addr));
    /* Send received datagram back to the client */
    if ((sendto(sock, echoBuffer, recvMsgSize, 0,
        (struct sockaddr *) &echoClntAddr,
        sizeof(echoClntAddr))) != recvMsgSize)
        printf("sendto() sent a different number of bytes
            than expected.\n");
}
}
```





# Run the Sample Programs (1)

---

- Give correct arguments

## Server process window

```
[shiyen@localhost 20071022]$ ./udpechosvr  
Usage: ./udpechosvr <UDP SERVER PORT>
```

## Client process window

```
[shiyen@localhost 20071022]$ ./udpechoclt  
Usage: ./udpechoclt <Server IP> <Echo Word> [<Echo Port>]
```



# Run the Sample Programs (2)

- Use correct username

## Server process window

```
[shiyang@localhost 20071022]$ ./udpechosvr 7  
bind() failed.
```

*Note:* binding the port number less than 1024 requires root authority

## Client process window

```
[shiyang@localhost 20071022]$ ./udpechoclt 192.168.1.253 hello
```



# Run the Sample Programs (3)

---

- Successful running using root

## Server process window

```
[root@localhost 20071022]# ./udpechosvr 7  
Handling client 192.168.1.253
```

## Client process window

```
[root@localhost 20071022]# ./udpechoclt 192.168.1.253 hello  
Received: hello  
[root@localhost 20071022]#
```



# Run the Sample Programs (4)

---

- Successful running using other username

## Server process window

```
[shiyen@localhost 20071022]$ ./udpechosvr 1500  
Handling client 192.168.1.253
```

## Client process window

```
[shiyen@localhost 20071022]$ ./udpechoclt 192.168.1.253 hello 1500  
Received: hello  
[shiyen@localhost 20071022]$
```